

I. Project Introduction

1. Overview

1.1 Project Information

This subsection records the institutional and academic context of the project. Project Title: Lightweight Deep Learning for Robust Shrimp Disease Detection on Mobile Devices under Diverse Environmental Conditions. Project Code: AIP491_01. Group Name: CVio. Academic Program: Bachelor of Artificial Intelligence. Department: Department of Artificial Intelligence. Institution: FPT University - Campus Can Tho. Supervisor: Nguyen Dinh Vinh. Project Team: Tran Huu Nhan, Le Thi Huynh Nhu, and Nguyen Van Phong.

Table 1.1: Project information

Item	Information
English project name	Lightweight Deep Learning for Robust Shrimp Disease Detection on Mobile Devices under Diverse Environmental Conditions
Vietnamese project name	Phát hiện bệnh tôm bằng mô hình học sâu nhẹ trên thiết bị di động trong nhiều điều kiện môi trường khác nhau
Project code	AIP491_01
Group name	CVio
Academic program	Bachelor of Artificial Intelligence
Department	Department of Artificial Intelligence
Institution	FPT University - Campus Can Tho
Project type	AI-based four-class image classification, BG/WSSV instance segmentation, robustness and explainability evaluation, deployment-oriented lightweight-model research, native Android application development, and Firebase-backed data integration
Software type	Native Android mobile application with integrated artificial-intelligence and data-management functions
Target platform	Android application using a conditional two-model TensorFlow Lite/LiteRT cascade on the common CPU execution path, with Firebase data services
Main AI tasks	Four-class shrimp disease image classification and BG/WSSV disease-region instance segmentation, with Healthy retained as an empty-label segmentation negative
Main technologies	Python, PyTorch, TorchVision, Ultralytics, OpenCV, Jupyter/Kaggle Notebook, TensorFlow Lite/LiteRT, Kotlin, Android Studio, Firebase, Git, and GitHub

Table 1.1 establishes that the project is primarily an AI research project with a mobile application as its deployment layer. The final application packages the classifier and instance-segmentation model as one conditional Android inference cascade; the project outcome therefore combines classification, localization, robustness, explainability, model-efficiency, and operational deployment evidence rather than treating the mobile prototype as the sole research contribution.

Table 1.2: Supervisor information

Full Name	Email	Mobile	Title
Nguyen Dinh Vinh	vinhnd29@fe.edu.vn	0975894685	Lecturer / Supervisor

Table 1.3: Team member information

Full Name	Email	Mobile	Role
Tran Huu Nhan (CE181655)	nhanthce181655@fpt.edu.vn	0867950216	Project Leader
Le Thi Huynh Nhu (CA182007)	nhulthca182007@fpt.edu.vn	0376387478	Member
Nguyen Van Phong (CE192081)	phongnv.ce192081@gmail.com	0342324690	Member

Tables 1.2 and 1.3 preserve the official project-contact information. The supervisor provides research direction and milestone assessment, while the project leader coordinates schedule, technical integration, reproducibility, and final document delivery.

1.2 Project Overview

This project investigates lightweight computer-vision methods for preliminary shrimp disease screening from images through two connected AI tracks. The classification track predicts Healthy, Black Gill disease (BG), White Spot Syndrome Virus (WSSV), or the WSSV_BG co-infection category. The instance-segmentation track localizes visible BG and WSSV disease regions. Healthy remains an image-level classification category but is not a positive segmentation class; Healthy images carry no BG/WSSV polygons and are retained in segmentation training, validation, and testing as empty-label negatives. The final application implements these tracks as a conditional two-model Android cascade. The FP32 classifier executes first for every accepted image. A Healthy result above the application threshold, or any below-threshold output, terminates the workflow after classification. An above-threshold BG, WSSV, or WSSV_BG output dispatches the FP16 instance-segmentation model, which produces only BG and WSSV mask instances. This dispatch rule controls application execution and must not be interpreted as verified disease ground truth. The project began with an approved scope centered on four-class classification, robustness evaluation under noisy conditions, lightweight mobile-edge deployment, and optional explainable-AI analysis. Instance segmentation was not part of the initially approved scope. It was added formally after the first review in Session 3 because classification alone was considered insufficient for the expected research depth. Later sessions introduced grouped-specimen splitting, attention-based segmentation refinement, Fourier-domain analysis, external dataset evaluation, and full Android/Firebase integration.

For segmentation validity, specimen grouping is enforced where specimen identity can be recovered from compatible filenames. Additional unmatched-name images do not expose compatible specimen identifiers and therefore cannot be claimed as specimen-grouped. The expanded segmentation inventory is consequently interpreted with an explicit grouping boundary rather than as fully specimen-leakage-free. Later controlled extensions also evaluate attention placement and Fourier-domain behavior.

Operational deployment evidence includes a retained 21-image real-world acquisition set collected across the three members' phone cameras and runtime observations from three Android devices using the common application/model bundle. The images cover variation in illumination, viewpoint, orientation, background clutter, and blur/softness. One unusably dark image was excluded. Because these photographs have no veterinary or laboratory ground truth, they support an operational case study




Fig. 1.1. Evolution of the project scope.

rather than field disease-accuracy validation. Figure 1.1 summarizes the chronological expansion of the project. The leftmost stage represents the initially approved classification scope; the middle stages show the review-driven addition of instance segmentation and later methodological refinements; the final stage presents the integrated research and deployment scope.

The figure is important for scope interpretation because it prevents the later segmentation work from being misrepresented as part of the original proposal. It also shows that grouped splitting, attention improvement, and Fourier-domain analysis were controlled scope extensions introduced after the segmentation track had become official.

Figure 1.1 should be read as a scope-evolution diagram rather than a post-defense chronology. Formal submission and Council-revision milestones are documented in Chapter II. The figure's purpose here is to distinguish the initially approved classification scope from the later formal segmentation track and its controlled methodological extensions. Figure 1.2 organizes the technologies by their actual function in the project. AI development tools are separated from experimentation infrastructure, mobile deployment technologies, and collaboration services so that the role of each technology is explicit.

The technology stack supports a traceable path from dataset preparation and model experimentation to Android inference and research-asset management. Kaggle/Jupyter and the accepted dual NVIDIA Tesla T4 training environment support reproducible experimentation; PyTorch and Ultralytics support model development; TensorFlow Lite/LiteRT and Android provide the deployment path; and GitHub, Google Drive, Google Docs, Google Meet, Firebase, and Zalo support controlled collaboration and application services. Training hardware, notebook/GPU timing, and Android application runtime are treated as separate provenance categories and are not numerically conflated.




Fig. 1.2. Project technology stack.

2. Project Background

Shrimp aquaculture is economically important, while infectious disease can disrupt production and farm operations. White spot disease is formally associated with infection by white spot syndrome virus (WSSV), for which aquatic-animal-health standards provide dedicated diagnostic methods. Image-based computer vision can therefore support preliminary visual screening, but it is not equivalent to veterinary or laboratory confirmation and must be interpreted within that boundary. can reduce irrelevant cues for classification, but it is not automatically beneficial for segmentation because contextual boundaries and local disease regions may be altered or removed. The project also addresses evaluation validity. Repeated views of one physical shrimp can produce optimistic localization estimates if random image-level splitting places related views in both training and testing partitions. Specimen grouping is therefore used where specimen identity is recoverable, while unmatched-name additions are explicitly separated because their filenames do not support specimen-level grouping. This distinction prevents the expanded inventory from being overstated as fully leakage-free. Mobile deployment provides an operational path for local preliminary screening. The final Android implementation uses a conditional classifier-to-segmentation cascade rather than executing both models for every image. Consequently, application runtime depends on which branch is executed; operational phone measurements are not automatically equivalent to isolated neural-network latency or to a controlled hardware benchmark.

3. Project Objective

The main objective is to design, evaluate, document, and deploy a lightweight deep-learning system for shrimp disease screening that combines class-balanced classification, disease-region localization,

robustness analysis, reproducibility, and mobile-oriented operational feasibility. The measurable objectives are as follows:

1. Organize and preprocess public shrimp disease image datasets using fixed and traceable data protocols.
2. Benchmark representative TIMM and official Ultralytics YOLO classification models and justify the selected deployment-oriented classifier under a preferred budget of approximately 15 million parameters or fewer; larger models may be retained as higher-capacity diagnostic controls rather than eligible final mobile candidates.
3. Improve the selected YOLO26m-cls backbone using ASL-LDAM and SimAM-DCFR, with Macro-F1 as the primary classification metric.
4. Evaluate the selected classifier on the clean test set, controlled synthetic corruption conditions, qualitative explainability visualizations, the separately trained EXT-3-Original regime, and the SDI-4 + EXT-3-Original partition-wise integrated retraining regime.
5. Create BG and WSSV disease-region masks for instance segmentation while retaining Healthy images as empty-label negative samples in training, validation, and testing.
6. Apply specimen-grouped splitting where specimen identity is recoverable, quantify leakage-prone image-level alternatives, and document the grouping boundary for unmatched-name additions whose filenames do not expose compatible specimen identities.
7. Evaluate attention-based and Fourier-domain segmentation improvements against controlled baselines.
8. Convert and integrate the final classification and conditional instance-segmentation artifacts into one native Android inference cascade using TensorFlow Lite/LiteRT, and assess branch-aware operational runtime on real Android devices.
9. Implement user and administrator workflows with authentication, camera/gallery input, prediction display, mask display, history, export/delete operations, and Firebase integration.
10. Maintain reproducible datasets, configurations, checkpoints, metrics, figures, and thesis evidence through controlled research-asset management.

4. Problem Statement

The project addresses the problem of developing a lightweight shrimp disease image-analysis system that balances class-balanced classification, disease-region instance localization, corruption robustness, repeated-specimen leakage control, model compactness, reproducibility, and mobile deployment feasibility. Strong performance on one clean split is insufficient if the model fails under acquisition disturbances, relies on spurious background cues, produces excessive disease predictions on Healthy negatives, or benefits from leakage between related specimen views. Conversely, model capacity alone cannot define lightweight suitability for an Android deployment target. The research problem is therefore multi-objective rather than a single-rank accuracy problem. Candidate selection must consider predictive quality together with parameter efficiency, negative-image behavior, protocol validity, and application-level execution constraints. The final mobile system adds a further systems consideration: classification always executes, whereas instance segmentation is conditional, so observed runtime depends on branch composition. Real-world three-camera photographs and three-device Android observations strengthen operational evidence, but the absence of veterinary/laboratory ground truth means they do not establish prospective field disease accuracy. Accordingly, the application is presented as preliminary image-based screening and research demonstration rather than clinical diagnosis. All conclusions remain bounded by the documented datasets, annotation protocol, split construction, corruption tests, and deployment measurements.

5. Significance of the Project

Scientifically, the project integrates four-class image classification and BG/WSSV disease-region instance segmentation within one controlled shrimp-disease study. It combines imbalance-aware classification, leakage-aware segmentation, corruption testing, qualitative explainability, attention-place-ment analysis, Fourier-domain diagnostics, and reproducible artifact management. This structure enables model-selection decisions to be supported by multiple evidence dimensions rather than by one aggregate metric. Practically, the project demonstrates a complete Android conditional cascade in which the classifier executes first and the segmentation model is invoked only for above-threshold disease outputs. The same application/model bundle was exercised across three Android devices, and the retained real-world image set was collected from three phone cameras under varied acquisition conditions. These observations demonstrate operational feasibility, but they do not prove farm-ready performance, clinical validity, universal device behavior, or iOS support.

6. Project Scope and Limitations

This section defines the work completed within the project and the boundaries that constrain interpretation. The initial proposal focused on four-class classification, robustness under noisy conditions, lightweight mobile-edge deployment, and optional explainability. Instance segmentation was later formalized as a required research track, followed by leakage-aware splitting, attention/Fourier investigations, additional-source classification experiments, and integrated Android/Firebase deployment.

6.1 Project Scope

The completed project contains two connected AI research tracks and one deployment layer. The classification track predicts Healthy, BG, WSSV, and WSSV_BG, with Macro-F1 used as the principal model-ranking metric. The instance-segmentation track predicts BG and WSSV disease-region masks. Healthy is not a segmentation mask class; Healthy images are retained as empty-label negatives. Classification external-evidence protocols distinguish three regimes. SDI-4 is the primary four-class classification regime. EXT-3-Original is a separately trained three-class source mapped to Healthy, BG, and WSSV. The SDI-4 + EXT-3-Original regime is a four-class integrated retraining experiment in which each source is split independently before train partitions are merged with train partitions, validation with validation, and test with test; WSSV_BG is supplied only by SDI-4. The integrated regime is therefore retraining rather than zero-shot external inference. The deployment scope integrates the final classifier and the conditional instance-segmentation model into one native Android application using the common TensorFlow Lite/LiteRT model bundle and Firebase-backed user/administrator workflows. The deployment study includes application-recorded branch-aware runtime observations on three Android devices and a 21-image three-camera operational acquisition set. Detailed experimental protocols remain in Chapter IV, and numerical outcomes remain in Chapter VI.

Table 1.4: Final project scope and explicit limitations

Scope component	Included work	Boundary / limitation
Classification	Predict Healthy, BG, WSSV, and WSSV BG from shrimp images; Macro-F1 is the principal ranking metric.	Evidence is limited to the documented datasets, source regimes, splits, and accepted checkpoints; it does not establish universal field generalization.

Continued on next page

Table 1.4: Final project scope and explicit limitations (continued)

Scope component	Included work	Boundary / limitation
Instance segmentation	Predict BG and WSSV disease-region masks; retain Healthy images as empty-label negatives.	Masks were produced by a single annotator without independent expert review. Specimen grouping is enforced only where identity is recoverable; unmatched-name additions cannot be claimed as specimen-grouped.
Robustness	Evaluate task-appropriate synthetic noise, blur, lighting, contrast, and related degradation conditions.	The finite corruption set does not represent every farm, camera, water condition, or real acquisition disturbance.
Explainability	Use heatmap-based classification XAI and diagnostic attention/Fourier analyses.	These analyses support qualitative or mechanistic interpretation under recorded protocols; they do not establish causal, clinical, or mask-accurate evidence.
Mobile application	Provide a native Android classifier-to-conditional-segmentation cascade with user/admin workflows and Firebase services.	Android only; no iOS validation. Application-recorded timing may include non-model overhead and is not guaranteed real-time performance on all devices.
Additional-source / integrated classification	Use EXT-3-Original as a separate three-class retraining regime and SDI-4 + EXT-3-Original as partition-wise integrated four-class retraining.	The integrated regime is not zero-shot external inference and is affected by source composition and class-prior differences.
Real-world operational evidence	Use 21 usable photographs collected across three member phone cameras and record 21 application runs per Android device.	Photographs lack veterinary/laboratory ground truth; image identity/order was not standardized across devices, so the study is operational runtime evidence rather than field accuracy or strict hardware ranking.
Intended use	Provide preliminary image-based screening, localized visual evidence, and reproducible research demonstration.	The system does not replace laboratory testing, veterinary assessment, treatment prescription, or autonomous clinical decisions.

Table 1.4 distinguishes completed project work from interpretations that the available evidence cannot support. In particular, it separates classification evidence from instance-segmentation validity, distinguishes Android operational runtime from pure model latency, and prevents three-camera photographs without expert ground truth from being interpreted as field disease-accuracy validation.

6.2 Project Limitations

Classification external validity remains bounded by the available source regimes. SDI-4 provides the primary four-class evidence. EXT-3-Original is evaluated as a separately trained three-class regime, while SDI-4 + EXT-3-Original is a partition-wise integrated retraining regime rather than direct external inference by the original SDI-4 checkpoint. Differences in source composition, acquisition conditions, and class priors therefore remain relevant to interpretation. Instance-segmentation validity is constrained by annotation and grouping information. BG and WSSV masks were produced by one annotator and were not independently reviewed by a second annotator or aquatic-animal-health expert. Specimen-grouped splitting reduces repeated-view leakage where filenames expose recoverable specimen identity, but the 303 unmatched-name additions do not provide compatible specimen identifiers; specimen-level separation cannot be enforced for that subset. Robustness tests cover selected synthetic disturbances and do not recreate every real farm, camera, lighting, water, disease-stage, or compression condition. Explainability is qualitative and does not establish causality or clinical validity. The three-camera operational photographs broaden visual conditions but have no veterinary/laboratory disease ground truth and therefore cannot support field accuracy claims. Deployment evidence is also bounded. The demonstrated implementation is Android-only and was exercised on three devices using the common application/model bundle. No iPhone/iOS runtime test was performed because a compatible iOS test device was not available. Each device contributed 21 application runs, but image identity and order were not standardized across devices because the experiment prioritized operational runtime feasibility rather than matched image-by-image hardware benchmarking. Application-recorded elapsed time may include preprocessing, interpreter activity, postprocessing, mask decoding, and UI-related work; it is therefore not automatically pure neural-network latency. The system is intended only for preliminary image-based screening and research demonstration. It does not replace laboratory testing, veterinary assessment, treatment prescription, or autonomous clinical decision-making. Future work should prioritize expert-labeled prospective field evaluation, multi-annotator segmentation review, controlled matched-image cross-device benchmarking, broader Android hardware coverage, and iOS deployment when suitable devices are available.

II. Project Management Plan

This chapter defines how the project was planned, coordinated, monitored, revised, and closed as an AI research project. The management plan is designed to support scientific validity as well as delivery discipline. It therefore covers not only schedule and task ownership, but also dataset integrity, experimental comparability, evidence traceability, mobile integration, document consistency, and the final submission constraint. The plan reflects the actual three-member working structure, the weekly supervision mechanism, and the transition from an initially classification-oriented scope to the final two-track classification and instance segmentation project.

The project differs from a conventional software-production project because the most important outputs are validated research evidence, reproducible model improvements, a defensible interpretation of results, and an integrated mobile demonstration. Progress is not measured only by completed code. A task is considered complete only when its inputs, configuration, output, evaluation, responsible owner, and relationship to the project objectives can be explained. This principle is applied throughout the team organization, WBS, risk controls, quality plan, change procedure, and Definition of Done presented in this chapter.

1. Team Work

1.1 Team Structure and Roles

The project team consists of one academic supervisor and three student members. The supervisor assigns weekly research tasks and reviews milestones. The project leader coordinates execution, deadlines, technical integration, reproducibility, and document delivery. After Session 3, responsibilities were separated into a classification track led by Tran Huu Nhan, an attention-based segmentation and mobile track led by Le Thi Huynh Nhu, and a dataset-labeling and Fourier-based segmentation track led by Nguyen Van Phong.

Table 2.1: Team structure and roles

Member	Student ID	Role	Main responsibilities
Nguyen Dinh Vinh	-	Supervisor	Weekly task assignment; research direction; milestone evaluation; academic feedback
Tran Huu Nhan	CE181655	Project Leader	Coordination; YOLO26m-cls selection; ASL-LDAM; SimAM-DCFR; classification evaluation; reproducibility; paper/thesis integration; TFLite support
Le Thi Huynh Nhu	CA182007	Member	Android requirements; UI/UX; Kotlin/Java; authentication; camera/gallery; TFLite; Firebase; admin dashboard; mask display; attention segmentation; documentation
Nguyen Van Phong	CE192081	Member	Initial preprocessing; baseline benchmarking; single-annotator BG/WSSV masks; YOLO conversion; grouped split; leakage analysis; segmentation baselines; Fourier experiments; robustness; documentation

Table 2.1 preserves the complete member information while clarifying the technical ownership of each workstream. It also shows that the supervisor is responsible for academic direction rather than implementation, whereas the project leader is responsible for coordination and integration across all tracks.



Fig. 2.1. Team structure and research responsibilities.

Figure 2.1 visualizes the reporting and responsibility structure described in Table 2.1. The hierarchy places the supervisor above the project leader, followed by the two specialized member tracks and the shared final-delivery responsibilities.

The figure is positioned directly within the Team Structure and Roles subsection because it is the visual representation of the organizational arrangement. The shared-responsibility node indicates that all members participate in final thesis review, formatting, presentation, and defense preparation.

Table 2.2: Roles and responsibilities

Role	Responsibility
Supervisor	Assign weekly research tasks; review evidence, scope, protocols, and milestones; provide academic approval and feedback.
Project Leader	Coordinate schedule and dependencies; support members; lead classification research; integrate results, mobile model support, reproducibility artifacts, paper writing, and thesis delivery.
Mobile and Attention Lead	Develop the Android/Firebase application and conduct attention-based instance segmentation research and documentation.
Dataset and Fourier Lead	Lead initial dataset analysis, baseline benchmarking, single-annotator mask creation, grouped splitting, leakage analysis, Fourier experiments, robustness evaluation, and technical documentation.
All Members	Review final thesis content, verify figures and tables, prepare presentation material, and support defense preparation.

1.1.1 Organizational Design Rationale

The organizational structure is deliberately small and specialization-oriented. A single supervisor provides academic direction and milestone evaluation, while the project leader connects the classification, segmentation, mobile, and documentation workstreams. This structure avoids creating artificial software-company roles that do not match the project. The team does not have separate product-owner, Scrum-master, quality-assurance, or database-administrator positions. Instead, responsibility is assigned according to the research artifact that must be produced and reviewed.

The project leader is positioned as the integration point because classification research, TensorFlow Lite conversion, reproducibility packaging, paper writing, and final thesis assembly share common dependencies. This role includes tracking whether datasets, class orders, checkpoints, figures, and reported metrics remain consistent across experiments and documents. Coordination therefore extends beyond deadline reminders: it includes detecting incompatible assumptions between workstreams and ensuring that accepted results can be incorporated into one coherent final project.

1.1.2 Role Evolution Across Project Phases

Role allocation changed when the research scope expanded. During the initial classification phase, Tran Huu Nhan identified the source dataset and Nguyen Van Phong led the first dataset inspection, pre-processing, and baseline-model benchmarking. After the first review and the formal addition of instance segmentation around Session 3, the work was separated more sharply. Nhan concentrated on YOLO26m-cls selection, ASL-LDAM, SimAM-DCFR, classification robustness, XAI, external classification experiments, and integration. Nhu combined Android/Firebase development with attention-based segmentation research. Phong took ownership of mask semantics, single-annotator labeling, grouped splitting, leakage analysis, segmentation baselines, and Fourier-domain experiments.

This transition allowed the team to preserve continuity in the classification track while adding a substantial localization track without abandoning the mobile application. It also created explicit integration points. Classification outputs and mobile class order had to remain synchronized; segmentation labels and mask displays had to use consistent BG/WSSV meanings; and both research tracks had to contribute traceable results to the same thesis. The team structure therefore reflects both specialization and controlled integration rather than isolated individual work.

1.1.3 Accountability, Review, and Shared Work

Primary ownership does not remove review responsibilities. The member marked as the operational owner must prepare the artifact, configuration, or result, while the project leader and supervisor verify whether the output can support the intended claim. For example, a completed experiment must identify the dataset split, seed, preprocessing, model configuration, checkpoint-selection rule, metric, and saved result. A completed annotation task must identify class semantics, format, negative-image handling, and the known limitation that no independent expert mask review was performed.

Shared responsibilities are concentrated near major delivery points. All members review the final thesis, inspect figure and table consistency, prepare the presentation, rehearse the demonstration path, and support defense preparation. This shared review is necessary because the thesis integrates multiple technical tracks and because a local inconsistency can affect the interpretation of the whole project. Collective review does not replace ownership; it functions as a final cross-workstream verification layer.

1.2 Communication Plan

Project communication is organized around one formal weekly Google Meet and continuous Zalo communication. The supervisor normally assigns tasks during the weekly meeting. The project leader then monitors deadlines, coordinates dependencies, and supports members. Most assigned work is expected before the next weekly review so that results can be discussed and used to define the next task cycle. Blockers are reported immediately through Zalo and escalated to Google Meet when direct technical discussion is required.

Table 2.3: Project communication plan

Communication item	Participants and purpose	Frequency	Tool / method
Weekly progress and task meeting	Supervisor and all members. Review completed work, inspect evidence, discuss blockers, assign next tasks and assign next tasks	Weekly	Google Meet
Technical blocker discussion	Relevant members; project leader; supervisor when required. Resolve data, model, mobile, Firebase, conversion, or integration blockers	As needed	Zalo and Google Meet
Experiment review	AI members and project leader. Check split policy, settings, checkpoints, metrics, plots, and interpretation	Weekly / as needed	Google Meet, Kaggle, Google Drive
Document review	Authors, reviewers, project leader, supervisor. Review thesis structure, wording, tables, figures, citations, and accepted revisions	Before milestones	Google Docs and Google Meet
Quick coordination	All members. Share status, short files, reminders, and urgent questions	Continuous	Zalo
Repository synchronization	Technical members. Synchronize source code, notebooks, configurations, and documentation	Before review and release	Git and GitHub

The communication plan does not claim formal Scrum events. Weekly meetings combine evidence review, feedback, replanning, and task assignment. Zalo is a supporting communication channel, while Google Drive, Google Docs, GitHub, and Kaggle contain the principal research and document artifacts.

1.2.1 Communication Cadence and Decision Flow

The weekly Google Meet is the formal decision point of the project. Each meeting begins with the results and unresolved issues from the previous task cycle. Members report what was attempted, which artifacts were produced, whether the expected acceptance condition was met, and which technical or scheduling constraints remain. The supervisor then provides direction and assigns the next work, while the project leader records ownership, dependencies, and the expected deadline. In most weeks, the deadline is placed before the following meeting so that the next discussion is based on current evidence rather than unverified progress statements.

Communication is therefore result-centered. A statement that training was completed is insufficient unless the relevant logs, checkpoint, metrics, and configuration are available. Similarly, a claim that a dataset is ready must be supported by counts, manifests, class mapping, mask format, or split checks. This evidence-centered cadence reduces ambiguity and makes weekly decisions easier to trace to the WBS and Meeting Minutes.

1.2.2 Blocker Escalation and Technical Resolution

Zalo is used for immediate coordination because technical blockers do not always occur near the scheduled weekly meeting. A blocker may concern dataset access, an invalid label, a failed training run, incompatible model operators, unexpected metric changes, Firebase behavior, document formatting, or dependency delays. The responsible member reports the issue as soon as it is identified, and the project leader determines whether it can be resolved directly or requires supervisor discussion.




Fig. 2.2. CRISP-DM applied to the project.

Escalation is proportionate to impact. A local implementation defect can be corrected without changing the project plan if it does not alter the dataset, metric, comparison protocol, or deadline. A blocker that changes a model configuration, invalidates a result, affects another member, or threatens a milestone is raised in Google Meet and recorded. This distinction prevents both extremes: silently changing important experimental conditions and over-formalizing minor technical corrections.

1.2.3 Communication Records and Artifact Traceability

Different communication tools serve different purposes. Google Meet supports formal review and task assignment. Zalo supports rapid coordination and short exchanges. Google Docs and Google Drive support collaborative writing, figures, reports, and meeting records. GitHub stores source code, configurations, and reproducibility documentation. Kaggle stores datasets, notebooks, training outputs, checkpoints, and result logs generated on the experiment platform. The project does not treat chat history as the final source of technical evidence.

A decision is considered traceable when the reason for the decision, the affected artifact, and the resulting update can be located. For example, the expansion to instance segmentation is reflected in Meeting Minutes, ownership assignments, WBS tasks, new dataset preparation, segmentation experiments, mobile mask-display requirements, and thesis scope. This multi-artifact consistency is important because management decisions and scientific claims must describe the same project state.

2. Project Management Approach

The project uses CRISP-DM as the named process model because it is an established data-mining and machine-learning lifecycle that matches the project emphasis on problem definition, data understanding, data preparation, modeling, evaluation, and deployment. The project does not rename CRISP-DM or claim a new official process model. Instead, the team applies the original phases iteratively and uses a weekly supervisor-controlled execution cycle to manage research tasks and evidence.

Figure 2.2 maps the original six CRISP-DM phases to the project. Dataset and experimental evidence are placed at the center because data quality, split validity, labeling, and traceable experiment outputs influence every phase.

The circular structure emphasizes iteration. Evaluation can send the project back to data preparation or modeling, and deployment findings can trigger further model or preprocessing revisions. Weekly




Fig. 2.3. Weekly research execution cycle.

supervision and evidence review operate across all phases rather than forming a seventh CRISP-DM phase.

Table 2.4: Application of CRISP-DM phases to the project

CRISP-DM phase	project application
Business Understanding	Define preliminary shrimp disease screening, four-class classification, robustness, model efficiency, and mobile deployment goals.
Data Understanding	Inspect ShrimpDiseaseImageBD and ShrimpDB, class distributions, image variability, duplicates, and repeated specimens.
Data Preparation	Perform classification-specific background removal, fixed splitting, mask annotation, class mapping, and grouped-specimen splitting.
Modeling	Benchmark TIMM and YOLO models; develop ASL-LDAM, SimAM-DCFR, attention segmentation, and Fourier-domain variants.
Evaluation	Use Macro-F1, Kappa, robustness testing, XAI, segmentation metrics, qualitative analysis, and leakage comparison.
Deployment	Deploy classifier + segmenter as a conditional Android TFLite cascade; integrate Firebase and user/admin workflows.

The formal CRISP-DM lifecycle is executed through the weekly research cycle shown in Figure 2.3. The supervisor assigns a task, members research and implement the method, experiments generate evidence, evaluation determines the result, and the weekly review decides whether to revise, continue, or replace the task.

This workflow is not a formal Scrum sprint because the project does not conduct daily Scrum events, formal sprint retrospectives, or a software-testing-driven increment for every week. It is a research execution pattern in which experimental evidence, implementation feasibility, and supervisor feedback determine the next action.

CRISP-DM was selected because it provides a recognized lifecycle for projects in which data characteristics, preparation choices, model behavior, evaluation findings, and deployment constraints repeatedly influence one another. The framework is more suitable than a formal Scrum description for the project because weekly work is driven by research questions and experimental outcomes rather than by a fixed software increment and a complete test suite. The six CRISP-DM phase names are retained without inventing a new methodology name.

Business Understanding establishes the initial four-class classification, robustness, lightweight deployment, intended users, and non-clinical boundary. Data Understanding investigates source datasets, class distribution, image variability, repeated specimens, and annotation feasibility. Data Preparation is task-specific: classification uses background removal and fixed image-level partitions, whereas segmentation requires team-created disease-region masks, Healthy empty-label negatives, and later a grouped-specimen split. These differences prevent the project from applying one preprocessing assumption indiscriminately to both tasks.

Modeling includes controlled screening of TIMM and official Ultralytics classification models, selection of YOLO26m-cls under the performance and approximate parameter constraint, implementation of ASL-LDAM and SimAM-DCFR, and the segmentation attention and Fourier studies. Evaluation checks clean performance, corruption behavior, classwise outcomes, qualitative XAI, mask metrics, leakage effects, and failure cases. Deployment converts selected outputs into Android, TensorFlow Lite, Firebase, farmer, and administrator workflows. Findings in any later phase may require revision of an earlier phase.

The weekly execution cycle operationalizes CRISP-DM without changing its official phases. A weekly task may begin in one phase and return to another after evaluation. For example, an unexpected segmentation result may trigger a label audit, a grouped-split revision, or a new modeling hypothesis. A mobile conversion failure may trigger architecture or preprocessing changes. This iterative interpretation makes the management approach consistent with the actual uncertainty of AI research.

2.1 WBS

The Work Breakdown Structure organizes the project into planning, dataset preparation, model development, mobile integration, documentation, and closure tasks. One man-day equals eight hours, and the estimates include both active work and scheduled experiment execution.

Classification uses the fixed seed-42 nominal 70/15/15 split. Segmentation retains Healthy images as empty-label negatives and uses a nominal 80/10/10 specimen-grouped split when filenames expose compatible specimen identities. The 303 unmatched-name additions lack recoverable specimen IDs and therefore use image-level stratified splitting; the expanded inventory is not claimed to be fully specimen-leakage-free. Modeling screens TIMM and official Ultralytics classifiers. YOLO26m-cls is retained under a deployment-oriented preference of approximately 15M parameters or fewer; larger candidates remain diagnostic references rather than eligible final mobile candidates. Subsequent work covers ASL-LDAM, SimAM-DCFR, and attention/Fourier segmentation. Deployment uses a conditional Android TFLite classifier-segmenter cascade.

Table 2.5: Detailed Work Breakdown Structure and estimated effort

#	WBS Item	Complexity	Est. effort (man-days)
1	Project planning and governance		
1.1	Confirm initial project scope and intended users	Simple	2
1.2	Establish weekly supervisor review cycle	Simple	2
1.3	Create communication and escalation channels	Simple	2
1.4	Create GitHub repository and branch structure	Medium	3

Continued on next page

Table 2.5: Detailed Work Breakdown Structure and estimated effort (continued)

#	WBS Item	Complexity	Est. effort (man-days)
1.5	Create Google Drive and Google Docs workspace	Simple	2
1.6	Prepare meeting-minute templates and tracking rules	Simple	2
1.7	Maintain milestone and deadline tracking	Medium	6
1.8	Review project scope after Session 3	Medium	2
1.9	Integrate instance segmentation into the official scope	Medium	3
2	Classification dataset preparation		
2.1	Locate and verify ShrimpDiseaseImageBD source	Simple	2
2.2	Inspect class folders and image counts	Medium	3
2.3	Check duplicates and repeated specimens	Complex	5
2.4	Implement background-removal preprocessing	Complex	6
2.5	Audit preprocessing outputs by class	Medium	3
2.6	Create fixed seed-42 train/validation/test split	Medium	4
2.7	Export split manifests and class statistics	Simple	2
2.8	Prepare EXT-3-Original class mapping	Medium	3
2.9	Prepare SDI-4 + EXT-3-Original partition-wise merge	Complex	5
2.10	Verify cross-partition source integrity	Complex	4
3	Classification baseline benchmarking		
3.1	Define lightweight parameter budget and metrics	Medium	3
3.2	Benchmark ConvNeXt and representative TIMM models	Complex	8
3.3	Benchmark YOLOv8 classification variants	Complex	5
3.4	Benchmark YOLO11 classification variants	Complex	5
3.5	Benchmark YOLO26 classification variants	Complex	5
3.6	Collect parameter count, size, latency, FPS, and metrics	Complex	5
3.7	Compare TIMM and YOLO efficiency-performance trade-offs	Complex	4
3.8	Select YOLO26m-cls as the improvement backbone	Medium	2
4	Proposed classification method		
4.1	Review imbalance-aware loss literature	Complex	4
4.2	Design ASL-LDAM objective	Complex	6
4.3	Implement ASL-LDAM in the training pipeline	Complex	7
4.4	Review SimAM and feature-recalibration literature	Complex	4
4.5	Implement SimAM gated residual integration	Complex	6
4.6	Design and implement DCFR texture recalibration	Complex	8
4.7	Integrate SimAM-DCFR with YOLO26m-cls	Complex	7
4.8	Run component-level ablation study	Complex	8
4.9	Select and freeze the final proposed checkpoint	Medium	3
4.10	Evaluate clean test performance	Medium	3
4.11	Evaluate corruption robustness	Complex	7
4.12	Generate Grad-CAM-based XAI evidence	Complex	5
4.13	Run EXT-3-Original classification experiment	Complex	5
4.14	Run SDI-4 + EXT-3-Original classification experiment	Complex	5
4.15	Package classification reproducibility artifacts	Complex	6
5	Instance-segmentation dataset preparation		
5.1	Define BG/WSSV disease-region mask semantics	Complex	4
5.2	Create single-annotator labeling protocol	Medium	3
5.3	Label BG disease regions	Complex	12
5.4	Label WSSV disease regions	Complex	14
5.5	Label co-infection images with both mask classes	Complex	8
5.6	Retain Healthy images as empty-label negatives	Simple	2
5.7	Convert masks to YOLO segmentation format	Complex	5
5.8	Generate annotation statistics and examples	Medium	3
5.9	Define specimen identifiers from filenames	Complex	4
5.10	Create grouped-specimen split	Complex	6
5.11	Audit specimen overlap and leakage	Complex	5
5.12	Prepare ShrimpDB segmentation annotations	Complex	8
6	Instance-segmentation modeling		
6.1	Benchmark YOLOv8-seg variants	Complex	6
6.2	Benchmark YOLO11-seg variants	Complex	6
6.3	Benchmark YOLO26-seg variants	Complex	6
6.4	Select grouped-split segmentation baseline	Medium	3
6.5	Review lightweight attention mechanisms	Complex	4
6.6	Implement attention modules in segmentation backbone	Complex	8
6.7	Run attention segmentation experiments	Complex	9

Continued on next page




Fig. 2.4. Project Gantt chart.

Table 2.5: Detailed Work Breakdown Structure and estimated effort (continued)

#	WBS Item	Complexity	Est. effort (man-days)
6.8	Document attention ablation results	Medium	4
6.9	Review Fourier-domain augmentation methods	Complex	5
6.10	Implement Fourier transformation variants	Complex	8
6.11	Run Fourier segmentation experiments	Complex	9
6.12	Evaluate segmentation under task-specific corruptions	Complex	7
6.13	Analyze qualitative predictions and failure cases	Complex	5
6.14	Package segmentation reproducibility artifacts	Complex	5
7	Mobile application and Firebase		
7.1	Analyze farmer and administrator requirements	Medium	4
7.2	Design Android UI/UX flows	Complex	7
7.3	Implement registration and authentication	Complex	7
7.4	Implement camera and gallery image input	Complex	6
7.5	Package final two-model TFLite bundle	Complex	7
7.6	Implement conditional Android CPU inference	Complex	8
7.7	Implement classification result display	Medium	5
7.8	Implement segmentation mask display	Complex	7
7.9	Implement history, export, and delete operations	Complex	7
7.10	Integrate Firebase data services	Complex	8
7.11	Implement administrator dashboard	Complex	8
7.12	Capture three-device runtime evidence	Medium	4
8	Documentation, publication, and closure		
8.1	Write classification paper draft	Complex	10
8.2	Write segmentation technical sections	Complex	8
8.3	Write project introduction and management plan	Complex	7
8.4	Integrate methodology and experimental results	Complex	10
8.5	Audit terminology, figures, and tables	Complex	6
8.6	Verify metric and checkpoint traceability	Complex	6
8.7	Prepare final GitHub documentation	Complex	5
8.8	Prepare presentation and defense materials	Complex	6
8.9	Conduct final thesis review and formatting	Complex	8
8.10	Submit initial defense DOCX package	Simple	2
	Total Estimated Effort (man-days)		487

Table 2.5 lists 90 executable tasks with a total estimated effort of 487 man-days. These values are planning estimates used to control workload, dependencies, and completion milestones.

Figure 2.4 presents the baseline schedule from 11 May to the initial defense DOCX submission on 5 August 2026. Milestones mark project start, segmentation extension, Session 11, integrated revision, and the initial submission freeze. The Gantt chart is a planning model rather than an exact timestamp record. Dates are reconstructed from Meeting Minutes 1-11, milestones, dependencies, and the initial pre-defense submission on 5 August 2026. The later Council No. 03 revision cycle is tracked separately under change management. Weekly review compares planned and actual state, identifies unfinished dependencies, and adjusts the next task allocation.

Table 2.6: Responsibility assignments

Responsibility	Tran Huu Nhan	Le Thi Huynh Nhu	Nguyen Van Phong	Nguyen Dinh Vinh
Project planning and weekly tracking	R	S	S	D
Initial dataset analysis and preprocessing	S	S	D	R
TIMM and YOLO baseline benchmarking	S	I	D	R
YOLO26m-clis selection and proposed classification	D	I	S	R
Classification robustness and XAI	D	I	S	R
BG/WSSV mask annotation and YOLO conversion	S	I	D	R
Grouped split and leakage analysis	S	I	D	R
Attention segmentation research	S	D	R	R
Fourier segmentation research	S	R	D	R
Android, Firebase, and administrator dashboard	S	D	I	R
Reproducibility packaging and GitHub documentation	D	S	S	R
Final thesis integration and formatting	D	D	D	R
Presentation and defense preparation	D	D	D	R

Responsibility codes are defined as D = Do, R = Review, S = Support, and I = Informed. The matrix identifies operational ownership without replacing the detailed role descriptions in Tables 2.1 and 2.2.

2.1.1 WBS Design Principles

The WBS decomposes the project by deliverable and research responsibility rather than by generic development activity. Parent work packages represent planning and governance, classification data, baseline benchmarking, proposed classification, segmentation data, segmentation modeling, mobile integration, and final documentation and closure. Child tasks are sufficiently specific to assign an owner and define an output, but they do not descend into trivial operations that would make the schedule unreadable.

Research tasks are separated when they have different acceptance conditions. Literature review, implementation, controlled training, ablation, robustness evaluation, external evaluation, qualitative interpretation, and reproducibility packaging are not treated as one combined task. This decomposition makes it possible to identify where a claim originated and whether the required supporting evidence was actually completed.

2.1.2 Effort Estimation Assumptions

One man-day is defined as eight hours for estimation. Because the project is a capstone rather than contracted employment, the estimate is used as a workload and scheduling measure, not as payroll accounting. Work may occur on any day of the week, and the team is evaluated by completion of assigned tasks rather than by a fixed office calendar. The schedule therefore records estimated effort and deadlines without claiming a conventional five-day employment pattern.

The user-approved estimation policy includes monitored GPU execution within the task effort because a long experiment occupies an assigned task, requires configuration, observation, checkpoint handling, failure recovery, and result extraction. This does not imply that unattended compute is equivalent to continuous human labor. It indicates that experiment duration consumes project time and affects de-

dependencies. The 487 man-day total is consequently a planning aggregate that combines active work and allocated experiment completion time.

2.1.3 Dependency and Parallel-Execution Logic

Dependencies protect scientific validity. Classification model comparison depends on a fixed dataset split and consistent metrics. Proposed-method experiments depend on the selected YOLO26m-cls baseline. Corruption and XAI analysis depend on a frozen checkpoint. Segmentation benchmarking depends on mask conversion and the grouped split. Attention and Fourier comparisons depend on the same baseline protocol. Mobile inference depends on a validated model conversion and stable class order. Documentation depends on results that have passed consistency checks.

Parallelism is used when dependencies permit it. Android UI and application infrastructure begin early while model research continues. Attention and Fourier segmentation are developed by different members after the common dataset preparation becomes available. Classification paper writing begins before all final thesis sections are integrated, while reproducibility packaging overlaps late experiments. This controlled parallelism shortens the project calendar without treating dependent outputs as interchangeable.

2.1.4 Interpretation of the Main Work Packages

The planning and governance package establishes the initial scope, collaboration environment, weekly review cycle, and scope-extension record. The classification-data package verifies the source dataset, performs ShrimpXNet-inspired background removal for the classification track, creates the seed-42 partitions, and prepares source-wise external datasets. The baseline package compares representative TIMM and official YOLO models so that YOLO26m-cls is selected from measured evidence rather than preference.

The proposed-classification package covers ASL-LDAM, SimAM-DCFR, component evidence, clean and corruption evaluation, XAI, external retraining, and reproducibility packaging. Segmentation packages cover mask semantics, Healthy empty-label handling, grouped splitting and leakage, baselines, attention, Fourier variants, robustness, and failure analysis. Mobile work demonstrates practical integration. The final package controls papers, thesis assembly, GitHub

2.1.5 Schedule Monitoring and Re-planning

Schedule control prioritizes research validity over superficial completion. An experiment is not marked complete merely to preserve a date if the configuration is inconsistent or the evidence is missing. Conversely, non-critical polishing can be reduced when it threatens a required dataset, experiment, integration, or submission milestone. The WBS, Gantt chart, responsibility matrix, and Meeting Minutes together provide complementary views of task scope, time, ownership, and decision history.

2.2 Risk Management

Risk management focuses on threats that can invalidate research conclusions, delay integration, or produce unsupported deployment claims. Each risk is associated with a preventive or corrective response.

Table 2.7: Project risks and response plans

#	Risk description	Impact / likelihood	Response plan
1	Repeated-specimen leakage	Impact: Critical; likelihood: Medium	Group by specimen and audit overlap where filenames expose IDs; report unmatched-name additions separately because specimen identity is unavailable.
2	Single-annotator mask inconsistency	Impact: High; likelihood: Medium	Use one documented labeling policy for consistency; disclose the absence of independent expert review as a limitation.
3	Class imbalance and co-infection overlap	Impact: High; likelihood: High	Use Macro-F1, classwise analysis, ASL-LDAM, confusion matrices, and controlled ablations.

Continued on next page

Table 2.7: Project risks and response plans (continued)

#	Risk description	Impact / likelihood	Response plan
4	Inconsistent baseline settings	Impact: Critical; likelihood: Medium	Fix dataset split, seed, preprocessing, checkpoint selection, metrics, class order, and evaluation scripts for direct comparisons.
5	Non-reproducible experimental result	Impact: High; likelihood: Medium	Archive configurations, environment metadata, checkpoints, logs, and exported result packages.
6	GPU quota or long experiment time	Impact: Medium; likelihood: High	Prioritize critical experiments, checkpoint runs, use staged debugging, and adjust deadlines through weekly review.
7	Mobile conversion or operator incompatibility	Impact: High; likelihood: Medium	Validate both TFLite artifacts, tensor shapes, class order, decoder, supported ops, and four-thread CPU path before deployment claims.
8	Prototype latency too high	Impact: High; likelihood: High	Report branch-aware Android runtime; separate no-seg and segmentation-triggered paths; avoid pure-latency, strict ranking, and real-time claims.
9	Firebase/network interruption	Impact: Medium; likelihood: Medium	Separate local inference from service failures; test authentication and persistence flows before demonstrations.
10	Scope creep after review feedback	Impact: High; likelihood: Medium	Apply the change-management process; assess schedule and dependencies; update WBS and meeting minutes.
11	Overclaiming diagnosis or generalization	Impact: Critical; likelihood: Medium	Describe the system as preliminary screening; state dataset, annotation, corruption, mobile, and external-validation limitations explicitly.
12	Final document inconsistency	Impact: High; likelihood: Medium	Review before 5 August, then complete the Council-mandated cross-document audit and revised submission before 12:00 on 15 August 2026.

The highest-priority risks are those that affect scientific validity: leakage, inconsistent settings, non-reproducible evidence, and overclaiming. Schedule and application risks remain important, but they

2.2.1 Risk Identification and Prioritization

Risk identification is performed from the perspective of research validity, delivery feasibility, and claim defensibility. A risk is considered critical when it can invalidate the comparison, create leakage, disconnect a reported metric from its checkpoint, or cause the thesis to make a claim unsupported by the available evidence. Operational inconvenience alone is not automatically critical unless it threatens a milestone or prevents completion of a required workstream.

Impact and possibility are assessed qualitatively because the project does not have sufficient historical data for calibrated probabilities. The classifications are used to focus attention, assign preventive controls, and determine escalation. High-impact risks are reviewed during weekly meetings, while immediate failures are reported through Zalo. The risk table records the agreed response rather than pretending that all uncertainty can be eliminated.

2.2.2 Scientific-Validity Risks

Repeated-specimen leakage can inflate segmentation performance when related shrimp views cross partitions. Grouped splitting and overlap audits control this for filename-compatible records. The 303 unmatched-name additions lack specimen IDs and use image-level splitting, so the expanded inventory is not claimed to be fully specimen-leakage-free.

Single-annotator masks improve internal labeling consistency but create uncertainty about clinical correctness and boundary subjectivity. The project addresses this by documenting mask semantics, using one protocol, treating Healthy images as empty-label negatives, and disclosing the absence of independent expert review. The limitation is not hidden by the use of precise polygon format or model metrics. Similar caution applies to XAI, which provides qualitative model-attention evidence but not causal or veterinary proof.

Inconsistent experimental settings are controlled because they can invalidate baseline comparisons more directly than a small metric difference. Dataset version, split, seed, preprocessing, augmentation, optimizer, checkpoint selection, class order, and evaluation code must be fixed or explicitly identified when comparing methods. The final classification result of 91.01% Macro-F1 is required to be reported

consistently in the final document according to the accepted project consolidation, and related tables, figures, and conclusions must use the same interpretation.

2.2.3 Resource, Integration, and Schedule Risks

Kaggle GPU availability, long training time, failed runs, and storage limits can delay research. The response is to debug with staged experiments, preserve checkpoints, prioritize the experiments necessary for the declared contributions, and revise deadlines through the weekly process. The project avoids presenting every attempted run as equally important; compute is allocated first to matched baselines, proposed methods, ablations, and required validation.

Mobile deployment introduces conversion and compatibility risks. Unsupported operators, changed tensor order, preprocessing mismatch, class-order mismatch, memory use, and CPU latency can make a high-performing training checkpoint unusable in the application. The team therefore

Mobile deployment introduces conversion and compatibility risks. Unsupported operators, changed tensor order, preprocessing mismatch, class-order mismatch, memory use, and CPU latency can make a high-performing training checkpoint unusable in the application. The team therefore validates TensorFlow Lite outputs against the source framework before packaging, preserves model metadata, and avoids converting application-recorded elapsed time into an unsupported real-time claim. Firebase and network failures are separated from local inference so that service interruptions do not redefine the model result.

2.2.4 Risk Ownership, Residual Risk, and Escalation

The operational owner of a workstream is responsible for detecting and reporting its risks, but the response may require multiple roles. Phong owns labeling, grouped-split, leakage, and Fourier risks; Nhu owns mobile, Firebase, UI, mask-display, and attention-integration risks; Nhan owns classification protocol, proposed-method, metric consistency, reproducibility, and final integration risks; the supervisor reviews decisions that affect scope, scientific interpretation, or milestone acceptance.

Residual risk remains after mitigation. The revision adds real-shrimp photographs from three member cameras and runtime logs from three Android devices using a common app/model bundle. These observations have no veterinary or laboratory ground truth, so they do not establish field accuracy or universal generalization. No iOS/iPhone runtime or formal farmer UAT was conducted; these limits remain explicit.

2.3 Quality Management

Quality management is evidence-based rather than limited to software test cases. The project applies quality controls at dataset, experiment, model, application, artifact, and document levels.

Table 2.8: Quality management plan

Quality area	Control activities	Acceptance output
Dataset integrity	Verify counts, class mapping, duplicate handling, specimen groups, split overlap, mask format, and representative samples.	Audited manifests, statistics, and visual examples
Classification experiments	Use fixed seed-42 protocol, consistent preprocessing, validation-selected checkpoint, Macro-F1, classwise metrics, corruption evaluation, and traceable logs.	Reproducible clean and corruption results
Segmentation experiments	Use documented mask semantics, grouped split, healthy negatives, mask mAP metrics, qualitative predictions, and leakage analysis.	Comparable grouped-split segmentation evidence

Continued on next page




Fig. 2.5. Document and experiment artifact flow.

Table 2.8: Quality management plan (continued)

Quality area	Control activities	Acceptance output
Model improvement	Require improvement over the relevant baseline on the declared primary metric and provide component-level evidence without mixing incompatible runs.	Supported proposed-method claim
Mobile application	Verify authentication, camera/ gallery input, preprocessing, TFLite inference, result display, mask display, history, export/delete, Firebase, and administrator functions.	Functional end-to-end prototype
Reproducibility	Store dataset/version references, settings, environment metadata, code, checkpoints, logs, metrics, and release packages.	Traceable research package
Document quality	Use consistent section numbering, table format, figure placement, captions, terminology, citations, metric values, and limitations.	Submission-ready final DOCX

Figure 2.6 shows how collaborative research assets are managed. Google Docs supports thesis drafting, Google Drive stores shared documents and figures, GitHub maintains source code and reproducibility materials, and Kaggle stores datasets, notebooks, experiments, checkpoints, and logs. Zalo is shown separately as a supporting communication channel.

The distinction between primary assets and communication is a quality-control requirement. Zalo may be used for quick file exchange, but final reproducibility evidence must be retained in structured repositories such as GitHub, Google Drive, or Kaggle rather than relying on chat history alone.

2.3.1 Dataset and Annotation Quality Assurance

Dataset quality begins with source verification, class-count inspection, file-format checks, duplicate and repeated-view analysis, and explicit partition manifests. Classification preprocessing is audited by class

so that background removal does not silently delete files or change class counts. External datasets are mapped with documented class semantics and split by source before partition-wise merging. These controls ensure that reported metrics can be related to known data rather than an undocumented folder state.

Segmentation quality requires explicit label controls. BG and WSSV are the only positive mask classes. Healthy has no BG/WSSV polygon but remains in train, validation, and test as an empty-label negative. The 1,149 filename-compatible images are audited by specimen group; the 303 unmatched-name additions cannot be grouped and are tracked as an image-level stratified extension. Polygon correspondence, class semantics, and the single-annotator limitation remain explicit quality controls.

2.3.2 Experimental and Model-Comparison Quality

Classification comparisons are governed by a fixed seed-42 protocol, consistent image partitions, documented preprocessing, and a common evaluation script. Macro-F1 is the primary ranking measure because the four-class problem is imbalanced and includes a co-infection class. Accuracy, classwise scores, Kappa, parameters, size, latency, FPS, corruption results, and XAI may support interpretation, but they do not replace the declared primary selection rule.

Direct segmentation baseline claims require the same grouped split, class mapping, mask format, image size, duration, seed, and evaluation settings. The CA-to-SimAM placement study is controlled across candidate insertion points. The broader attention-family comparison is not a strict one-factor ablation because archived runs differ in augmentation, epochs, software, initialization, placement, and auxiliary loss. Fourier variants use matched controls; qualitative panels supplement mask metrics with Healthy false positives, incomplete masks, merged regions, and difficult appearances. protocol is controlled, and the evidence supports the stated contribution. A metric increase alone does not establish novelty. The thesis must distinguish adopted components, team modifications, reimplementations, and newly proposed combinations. Unsupported causal interpretations and claims that exceed the dataset or protocol are removed during quality review.

2.3.3 Reproducibility and Artifact Quality

Reproducibility quality connects every publishable number to a dataset version, split rule, seed, environment, configuration, evaluation command, checkpoint, and exported result. GitHub provides the structured source and documentation layer; Kaggle preserves executable notebooks and run outputs; Google Drive retains shared figures, documents, and packaged evidence. File exchange through Zalo is temporary and cannot be the only location of an accepted result.

Artifact review also checks internal consistency. The metric shown in a table must match the associated figure and discussion. Model names must not change meaning between methodology and results. Healthy must be described as a classification class and an empty-label segmentation negative, not as a segmentation mask class. ShrimpXNet must remain a literature reference for classification, whereas the proposed classification method is built directly on the selected YOLO26m-cls engineering baseline.

2.3.4 Mobile and Document Quality

Mobile quality is assessed through Android flows and operational runtime, not production readiness. The common build always runs the FP32 classifier and conditionally dispatches the FP16 segmenter. Three Android devices provide branch-aware runtime under the same app/model configuration; image order is not standardized because the study targets runtime feasibility, not strict hardware ranking. The system remains a preliminary screening prototype. Document quality is controlled at structural, visual, terminological, and evidential levels. Chapter numbering follows the official AIP491 template. Tables use consistent Times New Roman typography, page-safe widths, peach headers, centered cell text, and explanatory prose. Figures appear in the relevant subsection with captions and interpretation. Cross-references, abbreviations, dataset names, model names, dates, limitations, and final metrics are checked before submission.

2.4 Budget and Funding (Optional)

No direct project budget was incurred. The team used public datasets, open-source frameworks, free Kaggle Notebook GPU allocations, existing Android devices, Firebase free-tier services, institutional resources, and no-cost collaboration tools. The table documents the resource model without inventing monetary expenditure.

Table 2.9: Project resources and funding status

Resource category	Resource / funding source	Recorded cost status
Datasets	ShrimpDiseaseImageBD and ShrimpDB public sources	No direct cost
Model development	Python, PyTorch, TorchVision, Ultralytics, OpenCV	Open source
GPU experiments	Kaggle Notebook; up to $2 \times$ NVIDIA Tesla T4 when available	Free allocation
Mobile development	Android Studio, Kotlin, Java interoperability, TensorFlow Lite	Open source / existing equipment
Application data services	Firebase free-tier services	No recorded direct cost
Version control and collaboration	GitHub, Google Drive, Google Docs, Google Meet, Zalo	No recorded direct cost
Testing devices	Existing personal Android devices	No dedicated purchase

2.4.1 Resource-Utilization Model

The absence of a direct monetary budget does not mean the project has no resource plan. The team relies on public datasets, open-source libraries, free Kaggle GPU allocation, existing Android devices, Firebase free-tier services, university supervision, and no-cost collaboration platforms. These resources define the feasible model scale, experiment count, deployment approach, and schedule. The approximate sub-15M-parameter selection constraint is consistent with the mobile-oriented objective and the available deployment environment.

Compute is managed as a scarce shared resource. Baseline screening is used to narrow the candidate set before expensive improvement experiments. Failed or exploratory runs are first debugged at smaller scale when possible. Checkpoints and logs are preserved to avoid unnecessary repetition. Late-stage compute is reserved for matched comparisons, ablations, robustness, external evaluation, and the final evidence package rather than unlimited architecture search.

2.4.2 Cost Control and Reporting Boundaries

The project does not invent estimated prices for tools that were used without recorded payment. Open-source software is recorded as open source, free-tier services as no recorded direct cost, and existing devices as existing equipment. This reporting boundary avoids presenting an artificial financial analysis that cannot be verified. It also distinguishes direct expenditure from the non-monetary cost represented by member effort and experiment duration.

Resource changes are reviewed when they affect scientific comparability or delivery. Moving from one GPU environment to another, changing image size to fit memory, reducing epoch count, or replacing a mobile model format may alter the result and therefore cannot be treated only as a budget decision. The effect on protocol, schedule, and thesis interpretation must be assessed through change management.




Fig. 2.6. Controlled change-management process.

2.4.3 Sustainability and Maintainability of Resources

The selected toolchain supports continuation after the capstone because it is based primarily on widely used open-source frameworks and accessible platforms. Source code and documentation can remain in GitHub, experiment notebooks can be rerun on compatible environments, and the Android application can be maintained without a proprietary inference runtime. Firebase free-tier dependence is documented so that future scaling or policy changes are not overlooked.

Sustainability also depends on preserving knowledge, not only files. The final repository and thesis should explain dataset preparation, class semantics, split construction, model selection, proposed modules, evaluation commands, mobile class order, known limitations, and artifact locations. A future maintainer should not need private chat history to determine how the accepted result was produced.

2.5 Change Management Process

A change may be triggered by council feedback, supervisor direction, experimental evidence, a technical blocker, dataset findings, deployment constraints, or document review. The team does not apply changes informally when they affect scope, comparison validity, dependencies, or reported conclusions. The impact is discussed, the decision is recorded, and accepted changes are reflected in tasks, dates, deliverables, experiments, and thesis wording.

Figure 2.5 presents the controlled path from a change trigger to documentation. Impact assessment considers scope, schedule, dependencies, and evaluation implications before the supervisor and team accept or revise the proposal.

Table 2.10: Change categories and control requirements

Change category	Assessment criterion	Required control
Minor implementation change	Does not alter dataset, split, metric, architecture claim, or deadline	Project leader coordinates and records the update
Experimental protocol change	Affects preprocessing, augmentation, seed, checkpoint selection, metric, or evaluation	Supervisor and team review before results are treated as comparable

Continued on next page

Table 2.10: Change categories and control requirements (continued)

Change category	Assessment criterion	Required control
Scope extension	Adds a formal research track, dataset, application function, or deliverable	Discuss impact; update WBS, schedule, ownership, and thesis scope
Schedule change	Moves a milestone or depends on unfinished GPU, labeling, or integration work	Update responsible owner and revised deadline during review
Document interpretation change	Changes the meaning of baseline, proposed result, limitation, or contribution	Verify evidence and apply consistently across methodology, results, discussion, and conclusion

2.5.1 Change Initiation and Classification

A change is initiated when new information makes the current plan incomplete, invalid, or inefficient. Triggers include council feedback, supervisor direction, experimental evidence, dataset findings, implementation failure, deployment constraints, dependency delays, or document review. The initiator states what must change, why the current state is insufficient, and which workstreams or conclusions may be affected.

Changes are classified by impact. A minor implementation correction does not alter the dataset, protocol, architecture claim, metric, or deadline. An experimental-protocol change affects comparability and requires review before new and old results are placed in the same table. A scope extension adds a formal track, dataset, function, or deliverable. A schedule change moves a milestone or dependency. A document-interpretation change alters the meaning of a baseline, proposed result, limitation, or contribution and therefore requires evidence verification.

2.5.2 Impact Assessment and Decision

Impact assessment examines scope, owner workload, dataset requirements, implementation effort, GPU time, dependencies, mobile effects, documentation changes, and the validity of previously produced evidence. The team determines whether earlier work remains usable, must be rerun, or needs to be relabeled. A proposed change is accepted only when its value and required work are understood; otherwise it is revised, postponed, or rejected.

Earlier review feedback triggered the instance-segmentation scope extension; responsibilities, WBS items, and thesis sections were updated accordingly. The later Council No. 03 post-defense revision is recorded separately in Section 2.5.5. The supervisor has the strongest role when a change affects research direction, contribution, evaluation protocol, or milestone acceptance. The project leader coordinates the operational update and checks cross-workstream effects. The responsible technical member estimates implementation and evidence requirements. This distribution prevents a local code change from silently becoming a new scientific claim.

2.5.3 Implementation, Verification, and Documentation

An accepted change is translated into revised tasks, owners, dates, dependencies, outputs, and acceptance conditions. Implementation then follows the same research sequence as other work: investigate, implement, experiment, evaluate, and review. Verification checks whether the change solved the initiating problem and whether it introduced a new inconsistency. The result is documented in Meeting Minutes, the WBS or schedule where relevant, structured artifact stores, and the thesis.

The addition of instance segmentation demonstrates the complete process. Council feedback identified insufficient research depth in the initial classification-only scope. The team accepted a formal segmentation track, defined mask-annotation work, assigned attention and mobile work to Nhu, assigned annotation, grouped split, leakage, and Fourier work to Phong, retained classification ownership with Nhan,

and updated the final system and thesis scope. Later grouped splitting further refined the segmentation protocol after repeated-specimen concerns were identified.

2.5.4 Cross-Document Consistency Control

A change is incomplete if only one paragraph is updated. Scope changes must be reflected in the introduction, objectives, methodology, results, project plan, figures, limitations, conclusion, WBS, Gantt chart, and role descriptions where applicable. Experimental changes must update configuration descriptions, table labels, metric interpretation, and reproducibility artifacts. Terminology changes must be propagated consistently to prevent conflicting automated or human readings of the thesis.

Before final acceptance, the project leader audits high-risk terms and values: YOLO26m-cls selection, ASL-LDAM/SimAM-DCFR naming, the locked 91.01% Macro-F1, Healthy segmentation semantics, split boundaries, attention/Fourier interpretation, two-model dispatch, runtime provenance, external-data limits, and non-clinical use. This links change management directly to document quality.

2.5.5 Post-Defense Council Revision Cycle

The 5 August 2026 DOCX package is retained as the initial pre-defense submission freeze. Council No. 03 reviewed the project on 9 August 2026 and required mandatory revisions before 12:00 on 15 August 2026. This post-defense cycle is treated as a controlled change extension rather than a retroactive rewrite of the historical WBS and Gantt chart. Accepted work includes additional component and comparator evidence, Fourier-spectrum analysis, final two-model mobile documentation, three-device Android runtime evidence, real-world three-camera operational evidence, and response-letter traceability.

2.6 Closure and Evaluation

Project closure uses the revised-submission gate rather than the initial defense freeze. The initial DOCX was submitted on 5 August 2026; Council No. 03 issued revision requirements on 9 August; and the controlled revision is due before 12:00 on 15 August 2026. Closure requires evidence, reproducibility, application functionality, document consistency, response-letter traceability, and archive completeness.

Table 2.11: Project closure and Definition of Done

Closure criterion	Definition of Done
Classification research complete	The final proposed classifier is fully specified; Macro-F1 is reported consistently as 91.01%; clean and corruption comparisons are traceable; component-evidence wording is internally consistent.
Segmentation research complete	BG/WSSV labels, Healthy empty-label handling, the specimen-grouped 1,149-image subset, the unmatched-name expansion boundary, baselines, attention/Fourier evidence, metrics, qualitative analysis, and limitations are documented.
Improvement criterion satisfied	Each claimed proposed method is compared against the relevant controlled baseline and is supported by the declared metric and evidence.
Research novelty documented	The thesis clearly identifies what is adopted, reimplemented, modified, or newly proposed; unsupported novelty claims are removed.

Continued on next page

Table 2.11: Project closure and Definition of Done (continued)

Closure criterion	Definition of Done
Reproducibility package complete	Datasets, split rules, seed, environment, configurations, code, checkpoints, logs, metrics, figures, and model metadata are archived.
Mobile prototype complete	One common Android application/model configuration executes the FP32 classifier for every accepted image and conditionally invokes the FP16 segmenter; branch-aware operational runtime from three Android devices is archived without claiming field accuracy or iOS support.
Paper and thesis complete	Classification paper material, segmentation documentation, tables, figures, references, limitations, and conclusions are integrated and reviewed.
Final quality audit complete	Section numbering, A4 page geometry, Times New Roman typography, table/figure placement, yellow revision highlighting, terminology, metrics, dates, limitations, and cross-document consistency pass the final rendered audit.
Defense preparation complete	Presentation content, demonstration path, role allocation, likely questions, and evidence references are prepared.
Final submission complete	The 5 August initial defense package and the Council-mandated revised package due before 12:00 on 15 August 2026 are both traceable; accepted artifacts are archived.

Table 2.11 defines closure by evidence, reproducibility, functional integration, and academic consistency. The 5 August package remains historical; closure requires the Council-mandated revision and final audit.

2.6.1 Closure Readiness Review

Closure readiness is assessed against the revised 15 August 2026 submission gate. Each workstream confirms that required outputs are available, accepted limitations are stated, unresolved placeholders are removed, and cross-workstream dependencies have been integrated. Classification, segmentation, mobile, documentation, reproducibility, response-letter evidence, and defense revisions are reviewed as connected components of one project; a missing artifact in one component can prevent closure even when another component has a strong metric. The readiness review distinguishes mandatory corrections from desirable extensions. Mandatory work includes the Council-requested evidence, controlled baseline/proposed comparisons, split and mask documentation, two-model mobile integration, traceable runtime reporting, template-compliant chapters, consistent figures and tables, references, limitations, and the revised DOCX/PDF package. Additional experiments are accepted only when they can be completed and integrated without destabilizing the required revision.

2.6.2 Final Scientific and Reproducibility Audit

The scientific audit checks whether every main claim is supported by the declared protocol and whether the scope of interpretation matches the evidence. It verifies the YOLO26m-cls deployment-oriented selection rationale, component evidence, multi-source retraining boundary, Healthy segmentation semantics, original-versus-expanded split validity, YOLO11n-seg selection trade-off, attention placement study, Fourier-spectrum analysis, three-device Android runtime interpretation, qualitative real-world

evidence, and external-validation limits. Unsupported real-time, clinical, causal, iOS, field-accuracy, or universal-generalization claims are removed. The reproducibility audit checks dataset references, split construction, seed, training environment, configurations, checkpoint identity, evaluation scripts, result files, figures, tables, model artifacts, application build information, and repository documentation. Training on dual NVIDIA Tesla T4, notebook/GPU efficiency measurements, and Android application-recorded runtime are treated as three separate provenance categories and are not numerically conflated.

2.6.3 Artifact Handover and Project Continuity

Project handover is organized around structured assets. GitHub retains maintainable code and reproducibility documentation; Kaggle retains permitted datasets, notebooks, checkpoints, and run outputs; Google Drive retains the integrated thesis, figures, meeting records, response materials, and packaged evidence. The Android application, final FP32 classification artifact, final FP16 segmentation artifact, Firebase configuration documentation, device/runtime evidence, and presentation material are retained with the final archive. Zalo remains communication history rather than the authoritative research repository. Continuity information includes the conditions under which the accepted evidence remains valid. Future teams must know that segmentation masks were created by one annotator, Healthy is an empty-label segmentation negative rather than a mask class, specimen grouping is enforceable only where compatible filename structure exposes specimen identity, external data remain limited, background removal is classification-specific, and mobile runtime is device- and branch-dependent. The real-world photographs do not carry veterinary or laboratory disease ground truth, and the revised study does not establish iOS support.

2.6.4 Lessons Learned and Final Evaluation

The main management lesson is that AI research requires continuous alignment between data, implementation, evaluation, deployment, and documentation. A model result cannot be managed independently from the split that produced it, the checkpoint that was selected, the evaluation code that interpreted it, or the claim that appears in the thesis. Weekly evidence review, explicit task ownership, structured artifact storage, and controlled scope changes were therefore more useful than adopting a nominal software process that the team did not actually practice. A second lesson is that scope expansion must be accompanied by resource and validity controls. Adding instance segmentation increased research depth but introduced annotation semantics, single-annotator limitations, repeated-specimen leakage risk, attention/Fourier comparison boundaries, conditional mobile execution, and additional documentation. The post-defense Council revision reinforced the same principle: new evidence is useful only when its provenance, comparison limits, and cross-document consequences are controlled. Final evaluation combines scientific, engineering, and academic criteria. The classification and segmentation contributions must be supported against their relevant baselines under declared protocols; the Android application must demonstrate the approved conditional two-model functions; runtime evidence must remain branch-aware and non-clinical; reproducibility artifacts must be traceable; and the final thesis must communicate results, limitations, dates, and scope consistently. The 5 August 2026 package is the initial defense submission, while successful completion of the Council-mandated revision before 12:00 on 15 August 2026 constitutes the revised project-closure gate.